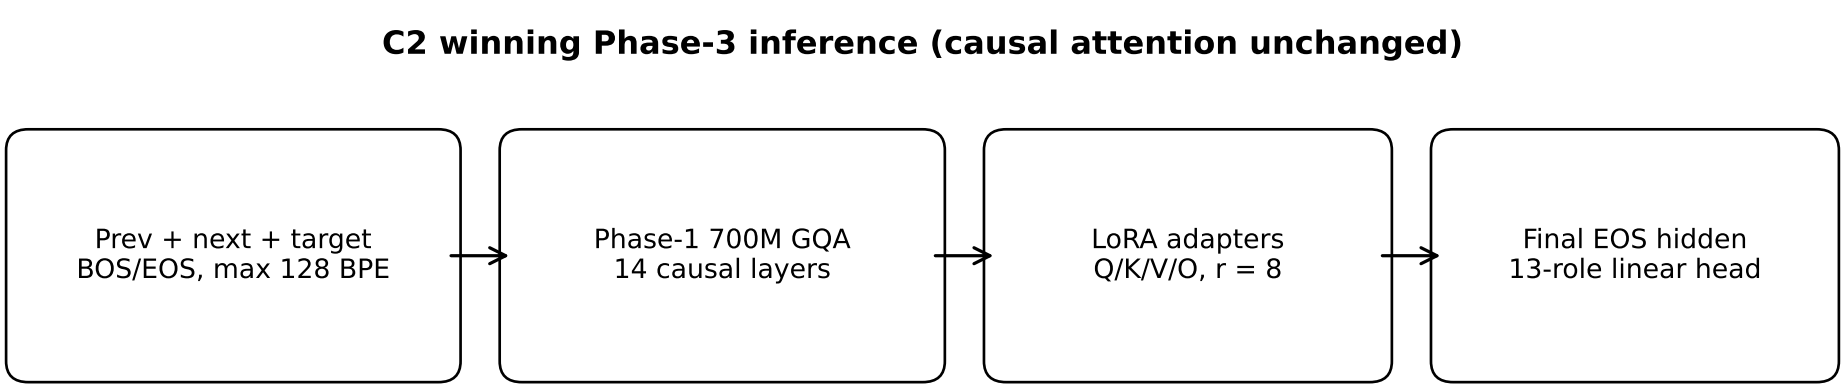
C2 winning Phase-3 inference (causal attention unchanged)
Prev + next + target
BOS/EOS, max 128 BPE
Phase-1 700M GQA
14 causal layers
LoRA adapters
Q/K/V/O, r = 8
Final EOS hidden
13-role linear head
Context is prepended in the token sequence; the target EOS can see both neighbors.